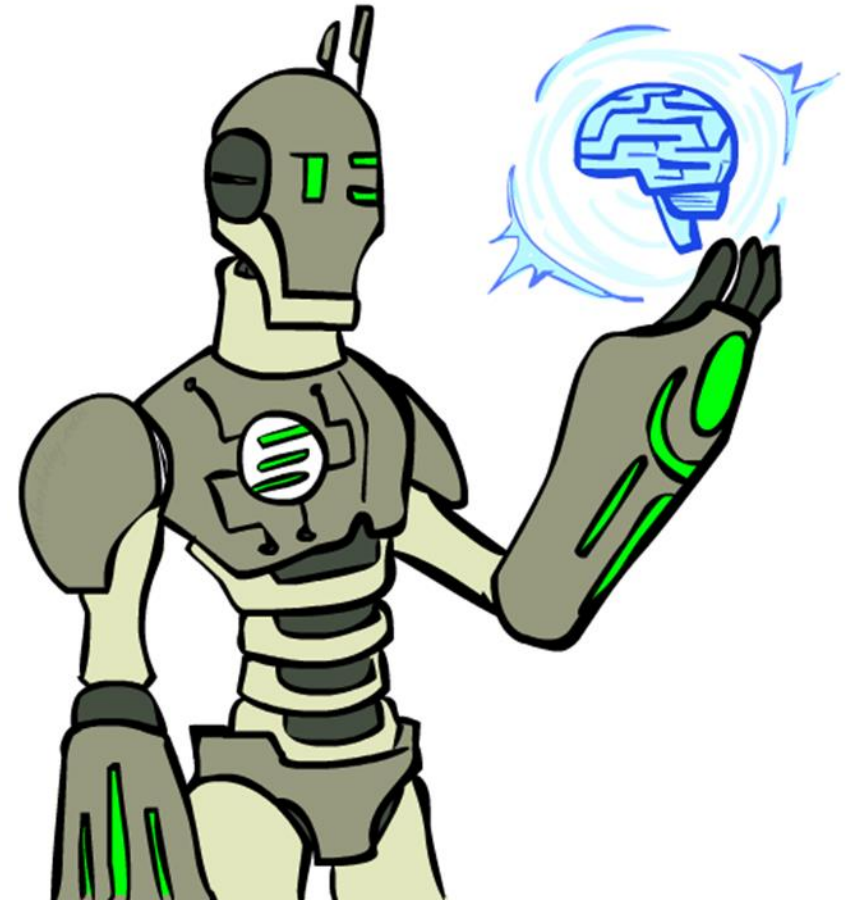


CISC440: Artificial Intelligence and Robotics

Chapter 1: Introduction

Outline

- ❖ What is Artificial Intelligence (AI)?
- ❖ The foundation of AI
- ❖ The history of AI
- ❖ The state of Art
- ❖ Risks and Benefits

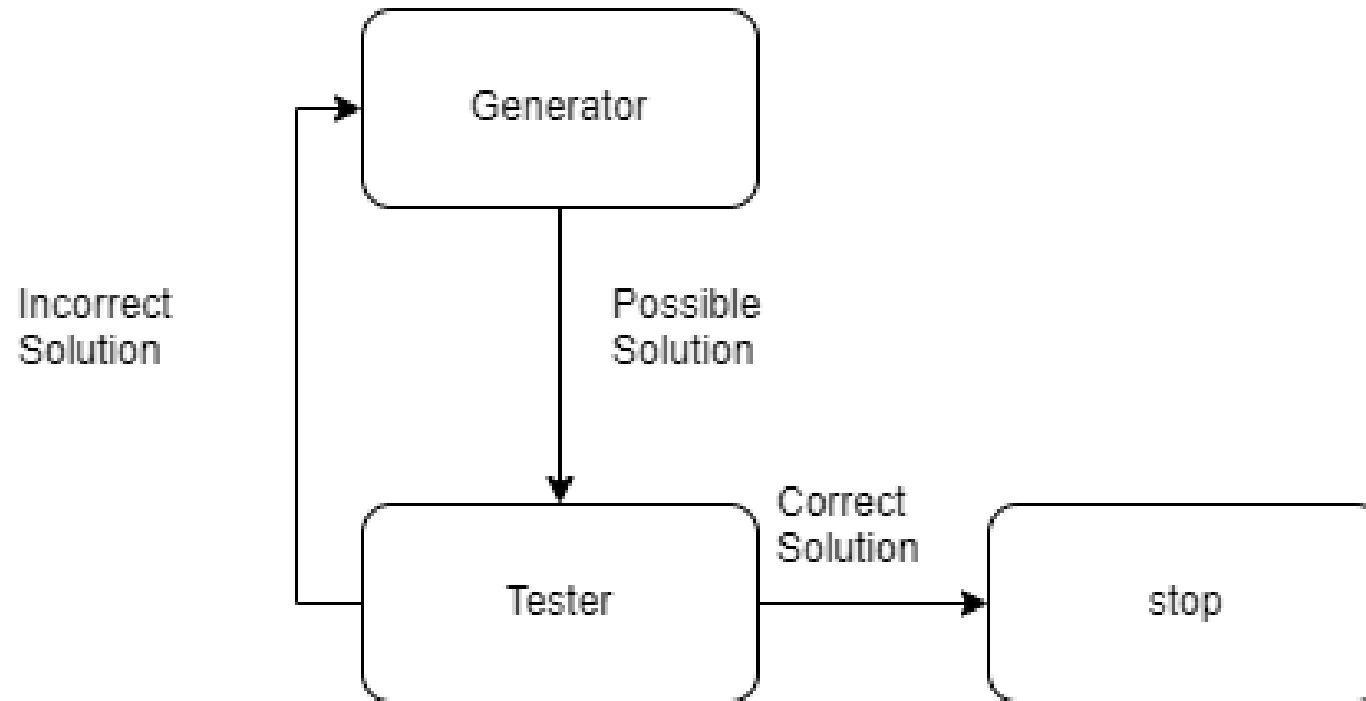


What is AI?

- ❖ Thinking + Perception + Action
- ❖ Philosophy class -> Fine
- ❖ Computer Science class ->
 - Models targeted at thinking, perception and action
 - Representation
 - Constraints exposed
 - Algorithm
- ❖ Example

**Algorithms enabled by constraint
exposed by representation that
support models targeted at thinking,
perception and action**

Generate and Test



How many continents cross the equator?



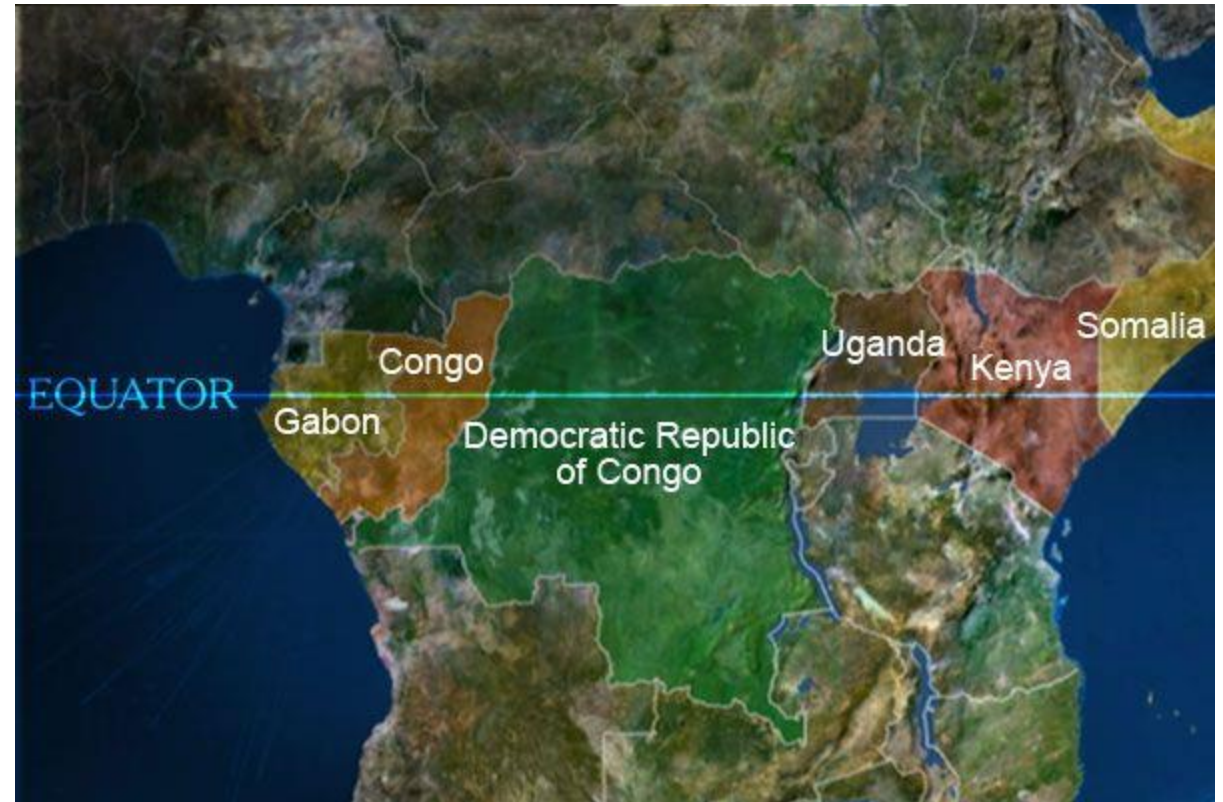
❖ The equator passes through the continents of South America, Africa and Asia.

How many countries of Africa cross the equator?

❖ Scan

❖ Interpret

❖ Answer



Sci-Fi AI?

Hope (70's)



80's Hardware could be scary



2000's can we be able to tell a difference



90's Software could be scary



Now: Reality Check

AI in News



Watson in Jeopardy

[ChatGPT](#)



Chatbots



AlphaGo



Autonomous Driving

AI at the Fun End

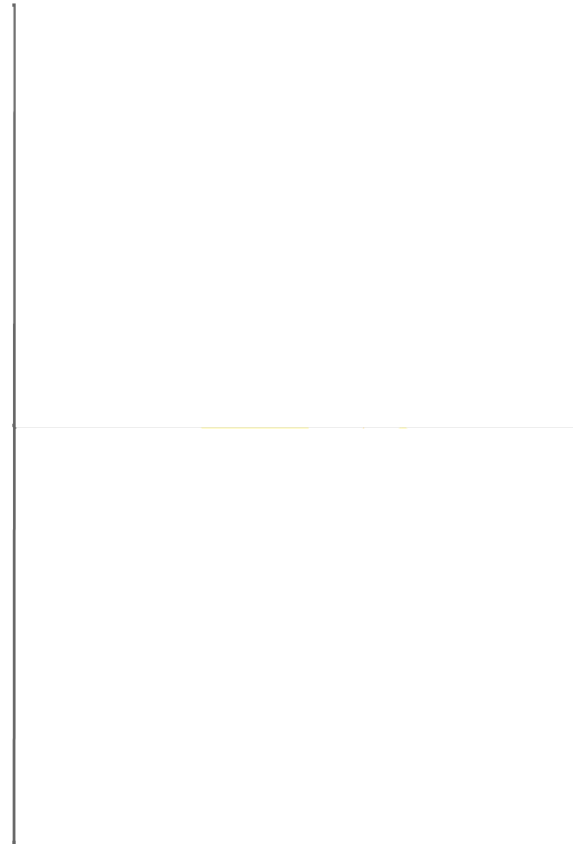
❖ Funny AI



Color Names

What is AI?

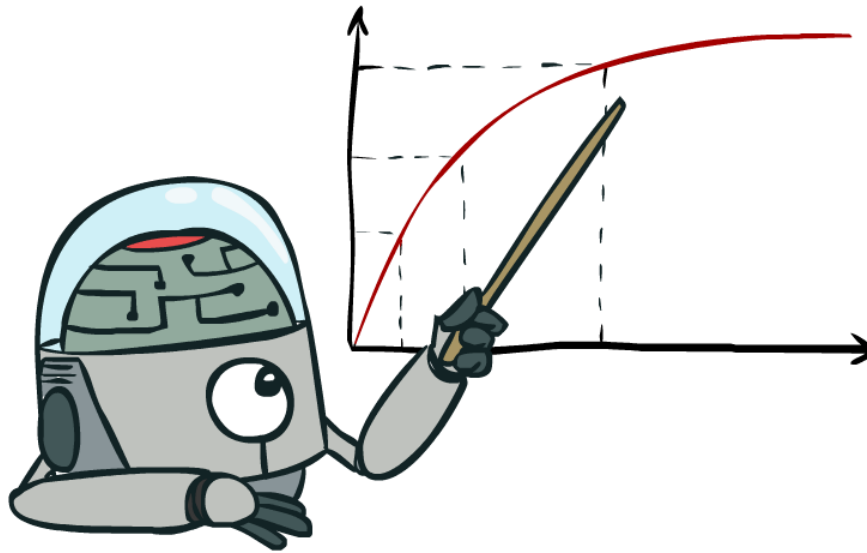
The science of making machines that:



Rational Decisions

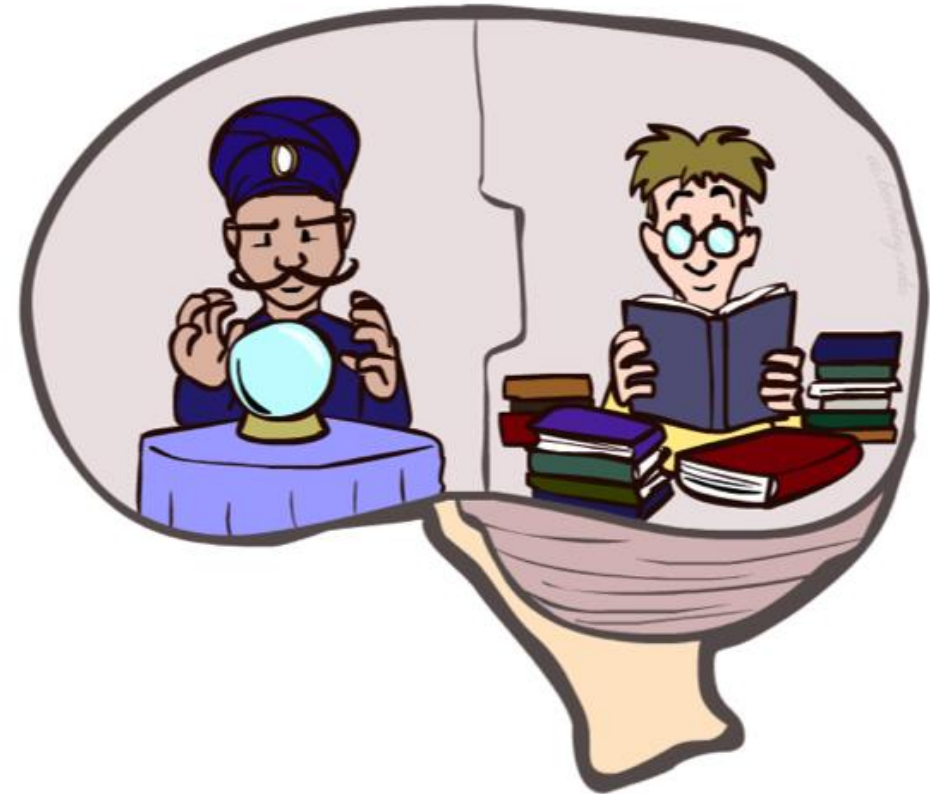
- ❖ We'll use the term **rational** in a very specific, technical way:
 - Rational: maximally achieving **pre-defined goals**
 - Rationality only concerns what decisions are made (not the thought process behind them)
 - Goals are expressed in terms of the **utility** of outcomes
 - **Being rational means maximizing your expected utility**

Maximize Your Expected Utility



What About the Brain?

- Brains (human minds) are very good at making rational decisions, but not perfect
- Brains aren't as modular as software, so hard to reverse engineer!
- “Brains are to intelligence as wings are to flight”
- Lessons learned from the brain: **memory and simulation** are key to decision making



Course Distribution

Part 1: Intelligence from Computation

Fast Search

Constraint Satisfaction

Uninformed Search

Part 2: Intelligence from Data

Bayes' Net

Decision Theory

Machine Learning

Part 3: Introduction of the Robotics

Connection

Motors

Intelligent Bricks

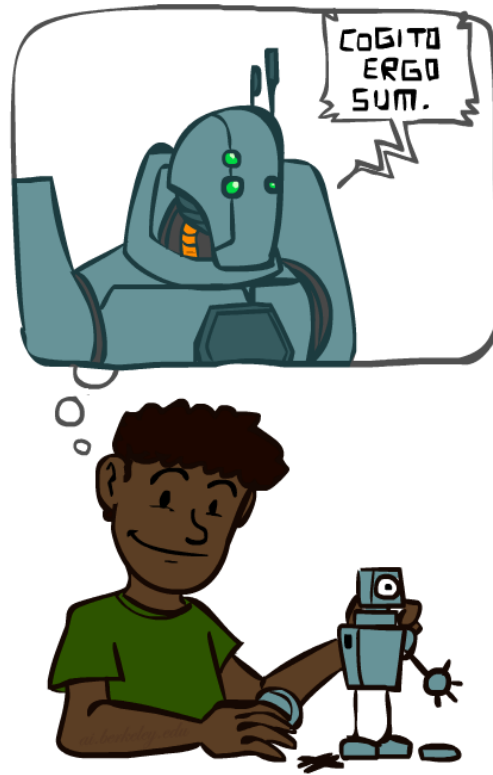
Sensors

LISP

Python

History of AI

Interesting interview from back in 1960's



History of AI

❖ 1842 Lady Lovelace:

- “The analytical engine has no presentation to originate anything. It can do whatever we know how to order it to perform”

❖ 1940-1950: Early days

- 1943: McCulloch & Pitts: Boolean circuit model of brain
- 1950: Turing's “Computing Machinery and Intelligence”

❖ 1950—70: Excitement- (a honeymoon period)

- 1950s: Early AI programs, including Samuel's checkers program, Newell & Simon's Logic Theorist, Gelernter's Geometry Engine
- 1956: Dartmouth meeting: “Artificial Intelligence” adopted
- 1965: Robinson's complete algorithm for logical reasoning

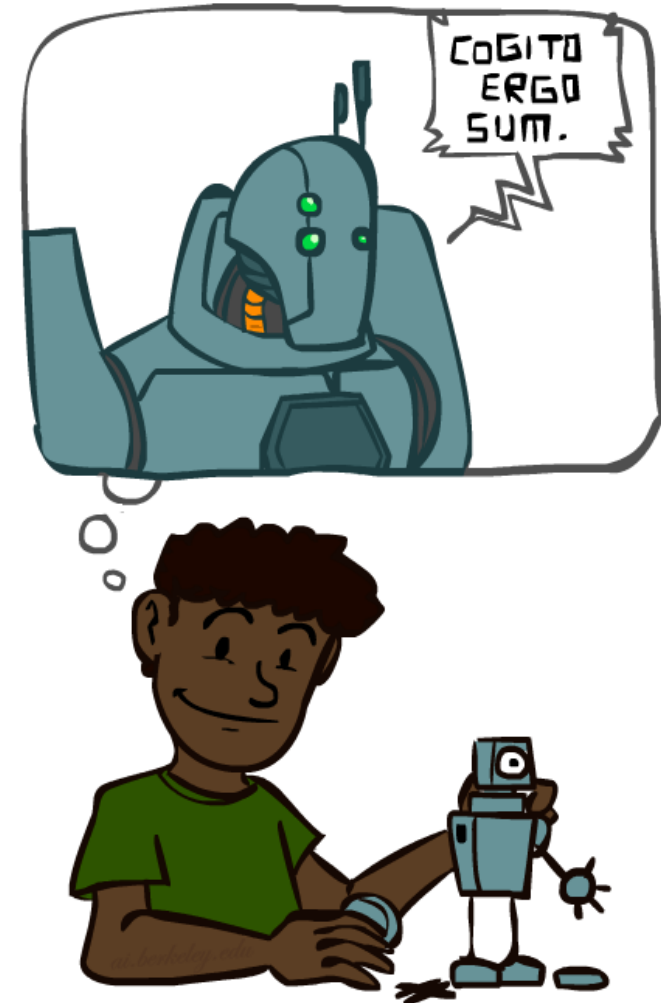
❖ 1970—90: Knowledge-based approaches (AI Winter)

- 1969—79: Early development of knowledge-based systems
- 1980—88: Expert systems industry booms
- 1988—93: Expert systems industry busts: “AI Winter”

❖ 1990—2000: Statistical approaches (AI Spring)

- Resurgence of probability, focus on uncertainty
- General increase in technical depth
- Agents and learning systems... “AI Spring”?

❖ 2000—: Where are we now?

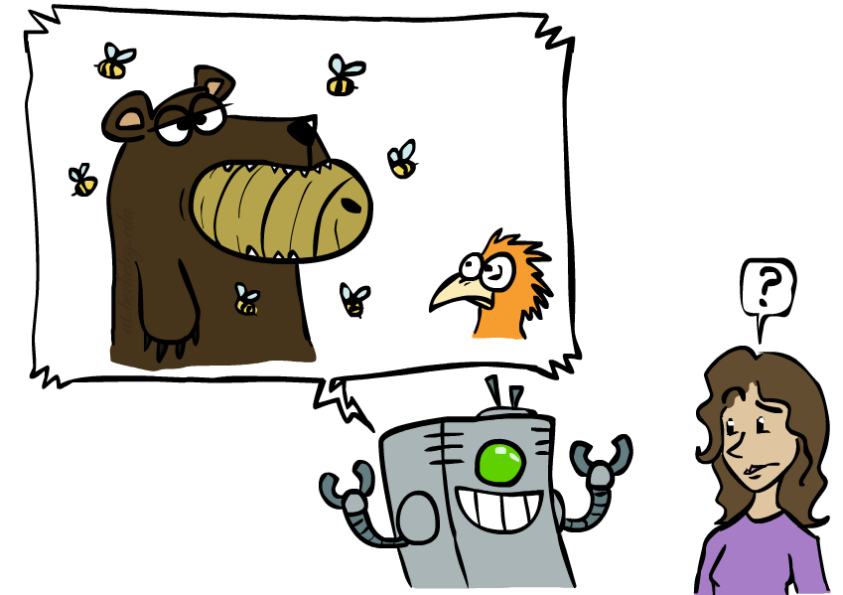


Unintentionally Funny Stories

❖ One day Joe Bear was hungry. He asked his friend Irving Bird where some honey was. Irving told him there was a beehive in the corner of the oak tree. He ate the bees.

❖ Henry Squirrel was thirsty. He went to the riverbank where his good friend was fishing. Henry slipped and fell in the water. The End.

❖ Once upon a time there was a vain crow. One day the crow was sitting in his tree, holding a piece of cheese in his mouth. He noticed that a fox was hungry, and swallowed the cheese. The fox walked over to the crow. The End.



ChatGPT

How does people think that Jan 1649 is an imaginary date?

Because they're buying stuff.

What were the skeletons in the closet.

They were bugs, mysterious.

Who is going to get the cookie?

Tiny one.

The most terrible of bears.

Once upon a time, there was a lazy dog named Bingo who lived in a small village. Bingo was so lazy that he wouldn't even bother to chase after his own tail.

One day, the village was having a big race and all the animals were excited to participate. Bingo, however, wasn't interested. He told his friends, "Why bother running when I can just sleep through the race?"

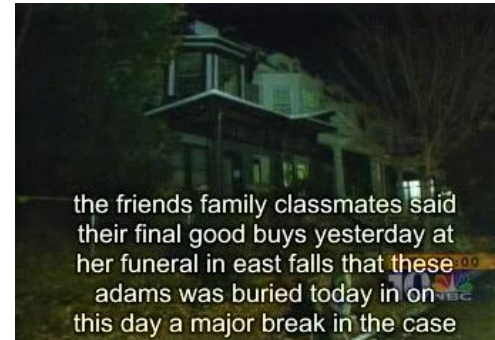
But when the race began, Bingo suddenly woke up from his nap and joined in. He was so fast that he zoomed past all the other animals, leaving them in the dust. When he crossed the finish line, he barked, "I didn't even know I was in a race!"

Bingo was declared the winner and everyone was shocked. From that day on, Bingo was known as the fastest dog in the village, but he still remained lazy and only ran when he was well rested.

Natural Language

❖ Speech technologies (e.g. Siri)

- Automatic speech recognition (ASR)
- Text-to-speech synthesis (TTS)
- Dialog systems



❖ Language processing technologies

- Question answering
- Machine translation

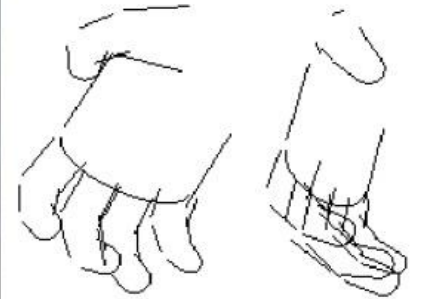
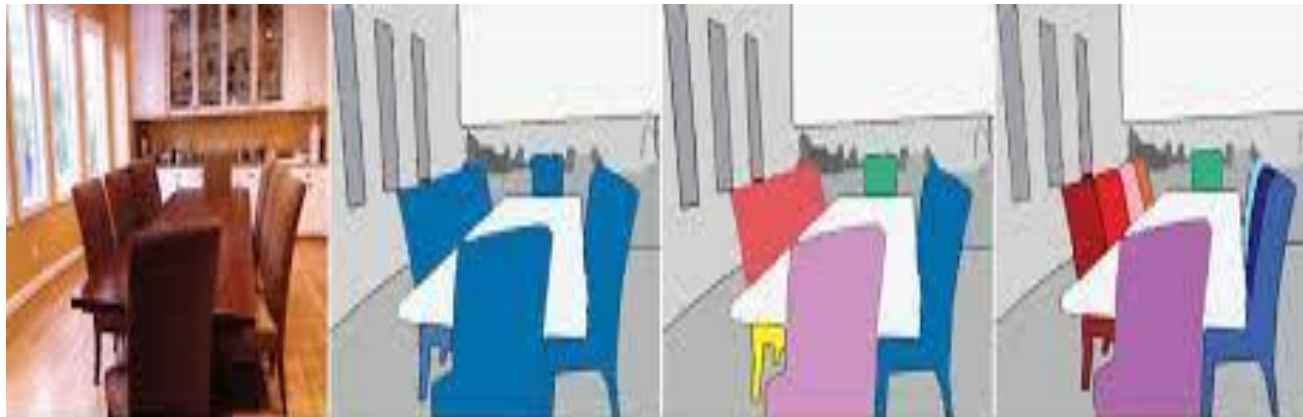
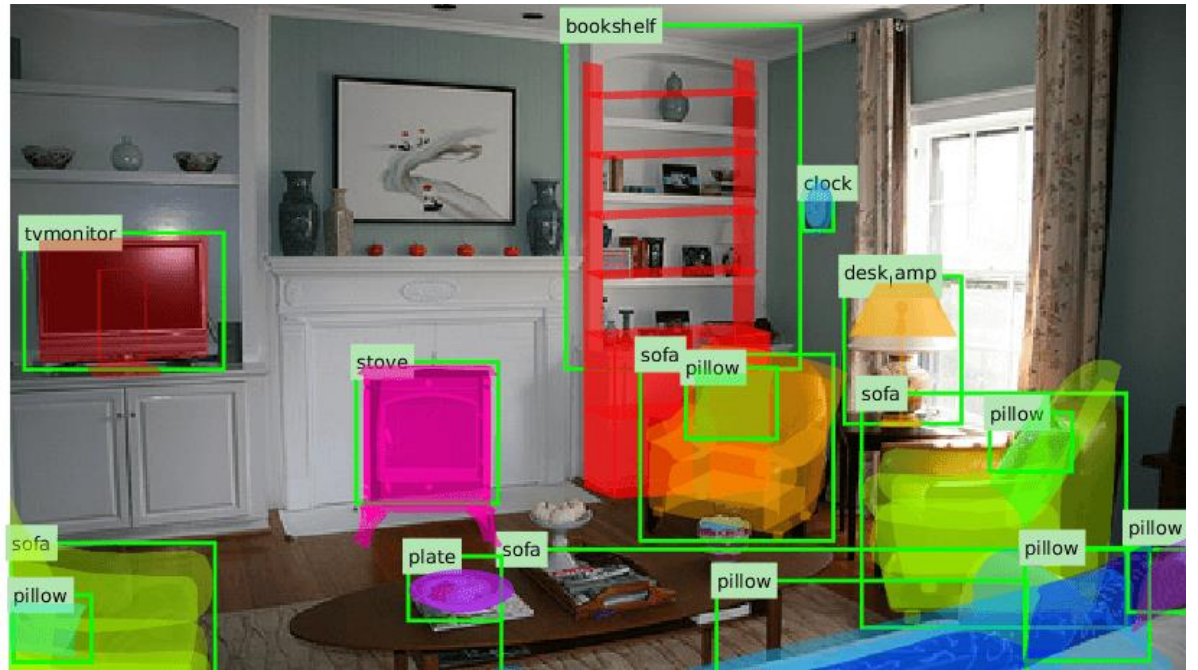


- Web search
- Text classification, spam filtering, etc...

Vision (Perception)

Pixel Value -> Info /decision

- Object and face recognition
- Scene segmentation
- Image classification



Computer Vision



"man in black shirt is playing guitar."



"construction worker in orange safety vest is working on road."



"two young girls are playing with lego toy."



"boy is doing backflip on wakeboard."



"girl in pink dress is jumping in air."



"black and white dog jumps over bar."



"young girl in pink shirt is swinging on swing."



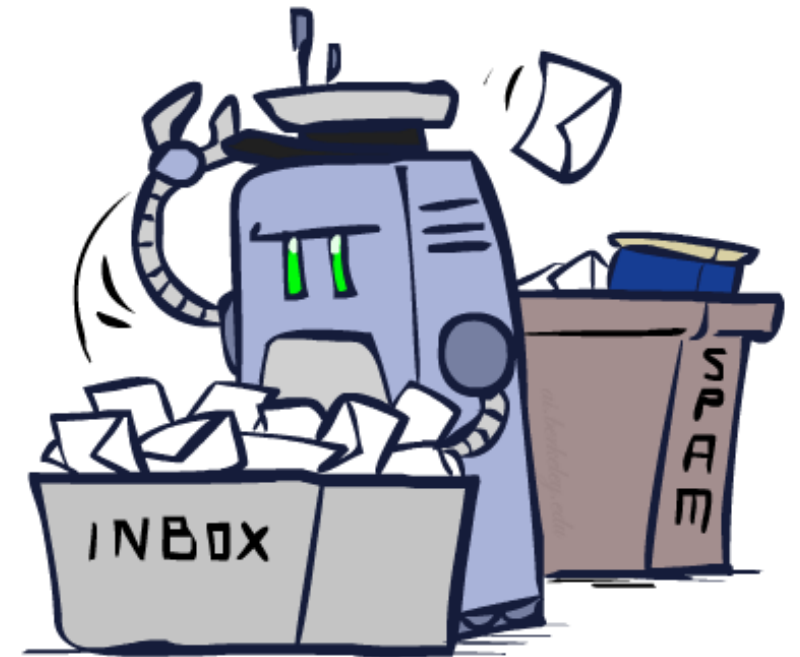
"man in blue wetsuit is surfing on wave."

Karpathy & FeiFei, 2015; Donahue et al., 2015; Xu et al, 2015; many more

Decision Making

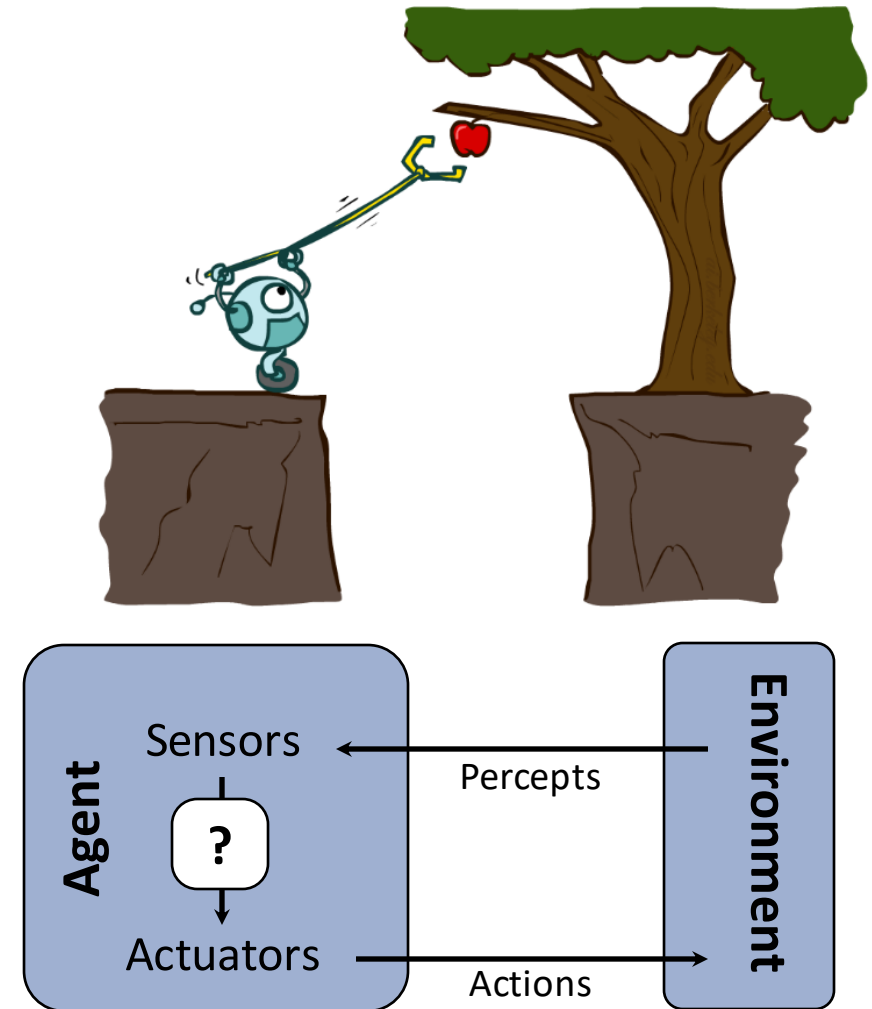
➤ Applied AI involves many kinds of automation

- Scheduling, e.g., airline routing, military
- Route planning, e.g., Google maps
- Medical diagnosis
- Web search engines
- Spam classifiers
- Automated help desks
- Fraud detection
- Product recommendations
- ... Lots more!



Designing Rational Agents

- ❖ An **agent** is an entity that *perceives* and *acts*.
- ❖ A **rational agent** selects actions that maximize its (expected) **utility**.
- ❖ Characteristics of the **percepts**, **environment**, and **action space** dictate techniques for selecting rational actions
- ❖ **This course** is about:
 - General AI techniques for a variety of problem types
 - Learning to recognize when and how a new problem can be solved with an existing technique



Summary

❖ Introduction

❖ History

❖ AI today

❖ Next time LISP Programming and more

Consent

Data for this course is taken from several books, slides and already similar course taught in other universities specifically MIT and AI-Berkley and is the sole property of the respective owner. The copyright infringement is not intended, and this material is solely used for the academic purpose or as a teaching material